

# Kubernetes / Raspberry Pi workshop Belastingdienst 24 nov 2016

## Environment config

Server	Adres
WLAN	Raspberry Pi Network / quintor2016
RPI SSH Login	host:rpi-node-* usr: root pwd: hypriot

## 1. Docker

Log in op je raspberry pi via ssh (root@<rpi-node hostname>) en voer de volgende Docker acties uit om basisvaardigheden met Docker op te doen. Als je niet bekend bent met Docker, is het sterk aangeraden dit door te lezen: <https://docs.docker.com/engine/understanding-docker/>. Als je later (dus niet nu) liever een video kijkt erover, dan hebben ze die ook, dus ga los: <https://training.docker.com/self-paced-training>.

1. Gebruik "docker help" en "docker <cmd> help" om informatie over de docker mogelijkheden te krijgen. Belangrijke commando's zijn onder andere `docker images` voor een lijst van de huidige images, `docker ps` voor een lijst van de huidige draaiende containers en `docker ps -a` voor een lijst van alle containers inclusief de inactieve.
2. Gebruik `docker run` om een container te starten van image `buildserver:5000/rpi-nginx` en maak deze beschikbaar op port 88. Controleer dat die draait door de webpagina op te vragen in je browser. Je hebt de flag `-d` nodig om de server als daemon te draaien (de command prompt blijft dan "hangen") en de flag `-p hostport:guestport` voor de poortmapping. Nginx serveert uiteraard op poort 80. Gebruik `--name` om je containerje een behapbaarder naam te geven dan het random ID. (hint: zorg dat de docker daemon draait, start hem anders met "docker daemon &"). Gebruik `curl` om de nginx webserver te benaderen "curl <http://localhost:88/>" vanaf de raspberry pi. Open de pagina ook vanuit je browser <http://<rpi-hostname>:88/>

Het geheel wordt dan iets als `docker run -d -p 80:80 --name mijn-nginx buildserver:5000/rpi-nginx`

3. Gebruik `docker exec -ti <container-id> bash` om een shell in de actieve container te openen. Bekijk met `ls` het bestandssysteem, met `ps -ef` de actieve processen, met `hostname` de hostnaam en met `hostname -i` het toegekende ip-adres van de container. (beëindig shell met `exit`)
4. Stop de docker container met `docker stop`. Check ook dat de webpagina niet meer beschikbaar is.
5. Met `docker ps -a` zie je welke docker container er zijn (actief EN niet actief).
6. Merk op dat elke keer dat je `docker run` uitvoert, er een nieuwe container aangemaakt wordt. Na een run kun je dezelfde container draaien met de toegewezen naam (`docker start name`) om een enorme wildgroei aan containers te voorkomen. Met de flag `-d` draait de container als een daemon en moet je hem expliciet stoppen. Start de container nu opnieuw.
7. Bekijk de lokaal beschikbare docker images met `docker images`.
8. Bekijk de image layering met `docker history`.
9. Verwijder de eerder gestarte docker container met `docker rm`. Gebruik de `-f` flag voor het geval de container nog niet gestopt is.
10. Verwijder de `buildserver:5000/rpi-nginx` image met `docker rmi`. Dit kan alleen als alle containers op basis van deze image verwijderd zijn. Als dit niet zo is, gebruik dan `docker ps -a` en `docker rm` om de rest te verwijderen.
11. We gaan nu iets spannenders doen, namelijk onze eigen (statische) webpagina serveren vanuit de nginx container. Download de statische pagina's "wget [http://buildserver/static\\_webpage\\_v1.zip](http://buildserver/static_webpage_v1.zip)"  
We hebben een nieuwe container nodig, dus we gebruiken `docker run`. Unzip `static_webpage_v1.zip` in een locatie naar keuze, bijvoorbeeld `/home/pirate/webpage1/`. "mkdir `/var/www/pirate`" "unzip `static_webpage_v2.zip` -d `/var/www/pirate/`".
12. Start een nieuwe container met een poortmapping en een mount naar `/usr/share/nginx/www/` met gebruik van `"-v hostfolder:guestfolder:mountoptions"`. Voor `mountoptions` kun je in dit geval `"ro"` gebruiken, voor read-only. Gebruik GEEN relatieve paden hierbij. Controleer met de browser de webpagina.

## 2. Kubernetes worker node

Op iedere Raspberry Pi is Kubernetes voorgeïnstalleerd om als worker node toegevoegd te worden aan het kubernetes cluster.

Let's check if everything is working correctly. Two docker daemon processes must be running.

```
$ ps -ef|grep docker

root      318      1  0 12:00 ?          00:00:24 /usr/bin/docker daemon -H
unix:///var/run/docker-bootstrap.sock -p /var/run/docker-bootstrap.pid
--storage-driver=overlay --storage-opt dm.basesize=10G --iptables=false
--ip-masq=false --bridge=none --graph=/var/lib/docker-bootstrap

root      697      1  8 12:01 ?          00:19:24 /usr/bin/docker daemon -H
fd:// --insecure-registry buildserver:5000 --storage-driver=overlay -D
--mtu=1472 --bip=10.1.62.1/24

root      30240 30106  0 15:48 pts/0      00:00:00 grep docker
```

The flannel container must be up.

```
$ docker -H unix:///var/run/docker-bootstrap.sock ps
```

CONTAINER ID	IMAGE	COMMAND
CREATED	STATUS	PORTS
		NAMES
2cceeaa7a06a	quay.io/coreos/flannel:v0.6.1-arm	"/opt/bin/flanneld
--" About an hour ago	Up About an hour	
kube_flannel_d1509		

The flannel network segment assigned to the node (see flannel0 => 10.1.xxx.0) must be used by the docker0 network bridge (10.1.xxx.1).

```
$ ifconfig

docker0  Link encap:Ethernet  HWaddr 02:42:2a:e1:bc:f2
          inet addr:10.1.62.1  Bcast:0.0.0.0  Mask:255.255.255.0
          inet6 addr: fe80::42:2aff:fe01:bef2/64 Scope:Link
          UP BROADCAST RUNNING MULTICAST  MTU:1472  Metric:1
          RX packets:45 errors:0 dropped:0 overruns:0 frame:0
          TX packets:43 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:0
          RX bytes:2800 (2.7 KiB)  TX bytes:6171 (6.0 KiB)

flannel0  Link encap:UNSPEC  HWaddr
00-00-00-00-00-00-00-00-00-00-00-00-00-00-00-00
          inet addr:10.1.62.0  P-t-P:10.1.62.0  Mask:255.255.0.0
          UP POINTOPOINT RUNNING NOARP MULTICAST  MTU:1472  Metric:1
          RX packets:0 errors:0 dropped:0 overruns:0 frame:0
          TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:500
          RX bytes:0 (0.0 B)  TX bytes:0 (0.0 B)
```

The hyperkube kubelet and proxy must be up.

```
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
cc27ddfd55a0	gcr.io/google_containers/hyperkube-arm:v1.3.6	"/hyperkube proxy --m"	About an hour ago	Up	About an hour
kube_proxy_5e947					
ef4ecc0c46da	gcr.io/google_containers/hyperkube-arm:v1.3.6	"/hyperkube kubelet -"	About an hour ago	Up	About an hour
kube_kubelet_17c52					

Once all the services on the worker node are up and running we can check that the node is added to the cluster on the master node. That can be done with `kubect1`, a tool for managing kubernetes. `Kubect1` for ARM can be downloaded from [googleapis storage](#). The default cluster server (master node) is configured, so that we don't have to specify it every time we run a command with `kubect1`. You can specify multiple clusters and use contexts to switch between them.

```
$ kubect1 config view
apiVersion: v1
clusters:
- cluster:
    server: http://rpi-master-1:8080
    name: rpi-cluster
contexts:
- context:
    cluster: rpi-cluster
    namespace: rpi-node-61
    user: ""
    name: rpi-cluster
current-context: rpi-cluster
kind: Config
preferences: {}
users: []
```

`kubect1 get nodes` shows which cluster nodes are registered along with its status.

```
$ kubect1 get nodes
```

NAME	STATUS	AGE
10.150.42.100	Ready	2h
10.150.42.101	Ready	2h

### 3. Kubernetes

## Running a container on kubernetes

An easy way to test the cluster is by running a simple docker image for ARM like the `rpi-nginx` one. `kubectl run` can be used to run the image as a container in a pod. `kubectl get pods` shows the pods that are registered along with its status. `kubectl describe` gives more detailed information on a specific resource. Each pod (container) get a unique ip-address assigned within the cluster and is accessible on that throughout the cluster, thanks to flannel overlay network. A pod can be deleted using `kubectl delete pod/deployment/service <name>`. Note that the `run` command creates a **deployment** (<http://kubernetes.io/docs/user-guide/deployments/>) which will ensure a crashed or deleted pod is restored. To remove your deployment, use `kubectl delete deployment nginx`. **(kubectl help is your friend!)** The `--port` flag exposes the pods to the internal network. The `--labels` flag ensures your pods are visible. Please use both flags.

Note that your nginx pod may seem to be stuck at "ContainerCreating" because it has to download the image first.

```
$ kubectl run nginx --image=buildserver:5000/rpi-nginx --port=80
--labels="run=nginx,visualize=true"
deployment "nginx" created

$ kubectl get pods -o wide
NAME                                READY    STATUS    RESTARTS   AGE      IP
NODE
nginx-1665122148-4amz1             1/1      Running   0           47s      10.1.87.2    10.150.42.100
$ kubectl describe pod nginx-1665122148-4amz1
Name:      nginx-1665122148-4amz1
Namespace: rpi-node-61
Node:      10.150.42.100/10.150.42.100
Start Time: Wed, 19 Oct 2016 02:18:38 +0200
Labels:    pod-template-hash=1665122148
           run=nginx
           visualize=true
Status:     Running
IP:         10.1.87.2
Controllers: ReplicaSet/nginx-1665122148
Containers:
$ kubectl get rs -o=wide
NAME                                DESIRED    CURRENT    AGE      CONTAINER(S)    IMAGE(S)
SELECTOR
nginx-1665122148                    1          1          3m      nginx            buildserver:5000/rpi-nginx
pod-template-hash=1665122148,run=nginx,visualize=true

$ kubectl get deployment -o=wide
NAME      DESIRED    CURRENT    UP-TO-DATE    AVAILABLE    AGE
nginx     1          1          1             1            4m
```

Now the container is running with kubernetes, the NGINX application is directly accessible via its IP address within the kubernetes cluster. Note that this is an IP address within the flannel overlaying network and is not accessible from outside the cluster. Also note that we do not have to specify any port mappings from the container to the host.

```
$ curl http://10.1.87.2/
<html>
<head>
<title>Welcome to nginx!</title>
</head>
<body bgcolor="white" text="black">
<center><h1>Welcome to nginx!</h1></center>
</body>
</html>
```

## Exposing containers on kubernetes

Now the pod is running, but the application is not generally accessible. That can be achieved by creating a service in kubernetes. The service will have a cluster IP-address assigned, which is the IP-address the service is available at within the cluster (10.0.0.\*). Use the IP-address of your Raspberry Pi node as external IP and the service becomes available outside of the cluster (e.g. 10.150.42.103 in my case). Check in your browser that `http://<ip-address-of-your-node>:90/` is available.

```
$ kubectl expose deployment nginx --port=90 --target-port=80
--external-ip=<my node's ip>
service "nginx" exposed

$ kubectl get svc
NAME          CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
nginx         10.0.0.102    10.150.42.103  90/TCP     57s
$ curl http://10.150.42.103:90
AND/OR (but not accessible to the outside world)
$ curl http://10.0.0.102:90
<html>
<head>
<title>Welcome to nginx!</title>
</head>
<body bgcolor="white" text="black">
<center><h1>Welcome to nginx!</h1></center>
</body>
</html>
```

## Scaling

The number of pod serving a service can easily be scaled via kubectl. Use `kubectl scale` to do so. Check out the visualizer the moment when you execute the scale command.

```
$ kubectl scale --replicas=3 deployment nginx
deployment "nginx" scaled

$ kubectl get pods -o=wide
```

NAME	READY	STATUS	RESTARTS	AGE
IP				
nginx-1665122148-4amz1	1/1	Running	0	13m
10.1.87.2				
nginx-1665122148-peep7	0/1	ContainerCreating	0	1m
<none>				
10.150.42.176				
nginx-1665122148-shm28	0/1	ContainerCreating	0	1m
<none>				
10.150.42.183				

## Doing a rolling update with Kubernetes (no service downtime)

In order to demonstrate a rolling update, we will use some prepared nginx containers which serve different static html depending on the version. Please remove your current deployment and deploy version 1 of this image, with 4 replicas, exposing port 80 on the pods (tip: if you don't remove the service exposing your raspi's port 90 to the world you can reuse it for this deployment).

```
$ kubectl delete deployment nginx
deployment "nginx" deleted

$ kubectl run nginx --image=buildserver:5000/rpi-nginx-withcontent:3
--port=80 --replicas=4 --labels="run=nginx,visualize=true"
deployment "nginx" created
```

Now we can see Kubernetes' full magic at work. We will edit the deployment (<http://kubernetes.io/docs/user-guide/deployments/#updating-a-deployment>) to start using the second version of the image, which will be rolled out by the system, replacing one pod at a time. The service will never go down, during the update a user simply gets served either the old or the new version. To kick the update off, you must edit the deployment. Take note of the different parts of this deployment file. You can write such a file yourself to deploy your applications, which is often more practical than having a bloke or gall hammer commands into a cluster with kubectl. For now, change the container image to version 4. For those unfamiliar with this editor, start editing with insert, stop editing with esc, save the result with :w and quit with :q. Alternatively, you can set the image directly.

```
$ kubectl edit deployment nginx
deployment nginx edited

OR

$ kubectl set image deployment/nginx
nginx=buildserver:5000/rpi-nginx-withcontent:4
deployment "nginx" image updated
```

Now let's assume that sometimes we inadvertently mess up and deploy a version of our application that is utterly broken. We get that dreaded midnight phonecall that a memory leak is destroying everything we care about, like uptime and service availability and professional pride.

Thanks to Kubernetes, we can run a single command from our laptop and get back to bed. First we will checkout the rollout history, pick a version to restore and then deploy it before snoring of happily.

```

$ kubectl rollout history deployment/nginx
deployments "nginx":
REVISION      CHANGE-CAUSE
1             <none>
2             <none>

# see details of a specific revision
$ kubectl rollout history deployment/nginx-deployment --revision=2
...

# execute the rollback
$ kubectl rollout undo deployment/nginx --to-revision=1
deployment "nginx" rolled back

```

For more details about the deployment rollback functionality, see <http://kubernetes.io/docs/user-guide/deployments/#rolling-back-a-deployment>.

## Creating services, deployments and pods from configuration files

Kubernetes resources can also be created from configuration files instead of via the command line. This makes it easy to put this Kubernetes configuration in version control and maintain it from there.

```

$ kubectl delete svc nginx
service "nginx" deleted

$ kubectl delete deployment nginx
deployment "nginx" deleted

$ kubectl create -f nginx-deployment.yaml
replicationcontroller "nginx" created

$ kubectl create -f nginx-svc.yaml
service "nginx" created

Now edit the deployment yaml file to use a different image version (4->3)
$ kubectl replace -f nginx-deployment.yaml

```

## 4. Deploying a three tier application

### Creating and claiming persisted volumes

The buildserver also hosts NFS service providing multiple volumes for mounting. In Kubernetes you can make a volume available for usage by creating a Persisted Volume.

**Edit the `nfs-pv.yaml` file so that the `nfs` share path matches your node. Also change the PV name to a unique value.**

```
$ kubectl create -f nfs-pv.yaml
persistentvolume "nfs-share-61" created
$ kubectl get pv
```

NAME	CAPACITY	ACCESSMODES	STATUS	CLAIM
nfs-share-61	1Gi	RWO	Available	

```
28s
```

Before the volume can be used it needs to be claimed for a certain application. This is done by creating a Persisted Volume Claim.

```
$ kubectl create -f nfs-pvc.yaml
persistentvolumeclaim "mysql-pv-claim" created

$ kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESSMODES	AGE
mysql-pv-claim	Bound	nfs-share-61	1Gi	RWO	12s

```
$ kubectl get pv
```

NAME	CAPACITY	ACCESSMODES	STATUS	CLAIM
nfs-share-61	1Gi	RWO	Bound	
rpi-node-61/mysql-pv-claim			5m	

More information can be found here: <http://kubernetes.io/docs/user-guide/persistent-volumes/>

Some of the deployment and service yaml files in the "assignment-3" folder are incomplete. Open the files and lookup the missing values in the Kubernetes documentation <http://kubernetes.io/docs/>.

- For the MySQL service, set an appropriate service type. Take some time to look at <http://kubernetes.io/docs/user-guide/services/#publishing-services---service-types> because services and their types are some of the most powerful and most important Kubernetes features.
- For the MySQL deployment, create Kubernetes secrets. Take a look at <http://kubernetes.io/docs/user-guide/secrets/#creating-a-secret-using-kubectl-create-secret> for more info. The MySQL root password is "root\_pw", the MySQL user is called "cddb\_quintor" and the MySQL password is "quintor\_pw". You can create secrets from the command line using:  
`kubectl create secret generic --from-literal=<field name>=<field value> <secret name>`  
so for example  
`kubectl create secret generic --from-literal=password=root_pw mysql-root-password`
- Choose a rollout strategy for the frontend deployment containers. Take a look at <http://kubernetes.io/docs/user-guide/deployments/#strategy>. It is not necessary to set the maxUnavailable and maxSurge fields but you can of course experiment with these values if you like.
- **For all three services (frontend, backend and MySQL), edit the \*-service.yaml files and set the IP of your own node before creating them.**

Now you can build the application from the ground up to the higher layers. Create the services and deployments for MySQL, backend and frontend:



```
$ kubectl create -f cddb-mysql-deployment.yaml
deployment "cddb-mysql" created

$ kubectl create -f cddb-mysql-service.yaml
service "cddb-mysql" created

$ kubectl create -f cddb-backend-deployment.yaml
deployment "cddb-backend" created

$ kubectl create -f cddb-backend-service.yaml
service "cddb-backend" created

$ kubectl create -f cddb-frontend-deployment.yaml
deployment "cddb-frontend" created

$ kubectl create -f cddb-backend-frontend.yaml
service "cddb-frontend" created
```

Test that the application is working using a browser and that it stores the data in the database.

You can scale up the frontend and backend layer. But the mysql layer cannot be scaled. Though Kubernetes manages the persisted volumes and remounts them on a different node when needed. To test this find out on which node the mysql pod is running and pull the network cable from the node and look what happens.